

NVIDIA NIM on Charmed Kubeflow Tutorial

This tutorial describes the complete steps to deploy a NIM (NVIDIA Inference Microservice) on Charmed Kubeflow from Canonical.

A NIM is a containerized inference microservice distributed from NVIDIA's private container registry - NGC (NVIDIA GPU Cloud).

By the end of this tutorial, you will be able to:

1. Deploy MicroK8s, a small, fast, single-package Kubernetes for datacenters and the edge.
2. Deploy Charmed Kubeflow, which enables developing and deploying machine learning models at any scale.
3. Deploy a NIM and serve a model with KServe, an addon component of Kubeflow.

This tutorial uses Kubeflow v1.8, Kubernetes v1.28, and Juju v3.1. Check the supported version combinations here:

<https://charmed-kubeflow.io/docs/supported-versions>

A GPU that is compatible with the model downloaded from NGC is required. Refer to the model details on NGC for further information. This tutorial is tested on a machine with an NVIDIA A100 GPU.

Deploy Ubuntu Server 22.04 LTS

~~# Deploy Ubuntu Server 24.04 LTS~~

<https://ubuntu.com/tutorials/install-ubuntu-server>

Update system

```
sudo apt update && sudo apt upgrade -y
```

Install NVIDIA GPU driver

```
sudo apt install nvidia-headless-535-server nvidia-utils-535-server -y
```

Reboot system and check NVIDIA GPU devices

```
sudo reboot
```

```
nvidia-smi
```

Install MicroK8s

```
sudo snap install microk8s --channel=1.28/stable --classic  
sudo snap install microk8s --channel=1.29/stable --classic
```

Add the current user to the microk8s group to avoid having to use sudo for every microk8s command

```
sudo usermod -a -G microk8s $USER  
  
newgrp microk8s
```

Enable MicroK8s add-ons needed to run Charmed Kubeflow

```
IP=$(hostname -I | awk '{print $1}')  
  
microk8s enable dns hostpath-storage ingress gpu metallb:"$IP-$IP"  
  
microk8s enable dns hostpath-storage nvidia ingress metallb:10.64.140.43-10.64.140.49
```

Check MicroK8s status until the output shows "microk8s is running" and the add-ons installed are listed under "enabled"

```
microk8s status --wait-ready
```

Add an alias for omitting microk8s when running kubectl commands

```
alias kubectl='microk8s kubectl'  
  
echo "alias kubectl='microk8s kubectl'" >> ~/.bash_aliases
```

Install Juju

```
sudo snap install juju --channel=3.1/stable  
sudo snap install juju --channel=3.5/stable
```

Configure MicroK8s to work properly with Juju

```
microk8s config | juju add-k8s microk8s-1 --client
```

Note: Command “microk8s config” retrieves the client’s Kubernetes config which is then registered to Juju Kubernetes endpoints.

Deploy Juju controller to MicroK8s

```
juju bootstrap microk8s-1
```

Add model for Kubeflow

```
juju add-model kubeflow
```

Deploy Charmed Kubeflow

```
juju deploy kubeflow --trust --channel=1.8/stable
```

```
juju deploy kubeflow --trust --channel=1.9/stable
```

Check juju status until all apps, except dex-auth, istio-pilot, and oidc-gatekeeper, become active

```
watch -c 'juju status --color | grep -E "blocked|error|maintenance|waiting|App|Unit"'
```

~~# Configure oidc-gatekeeper with the ingress-gateway IP~~

```
IP=$(microk8s kubectl -n kubeflow get svc istio-ingressgateway workload -o  
jsonpath='{.status.loadBalancer.ingress[0].ip}')
```

```
juju config oidc-gatekeeper public-url=http://$IP
```

~~# Check juju status until all apps become active~~

```
watch -c 'juju status --color | grep -E "blocked|error|maintenance|waiting|App|Unit"'
```

Create an account at <https://ngc.nvidia.com/signin> and create an API key at

<https://org.ngc.nvidia.com/setup/api-key>

Note: You must have NIM access on NGC.

Set an environment variable for the API key

```
export NGC_CLI_API_KEY=<key>
```

Create Kubernetes secret with the NGC API key to download NIMs from NGC private Docker registry

```
kubectcl create secret docker-registry ngc-docker-secret \  
--docker-server=nvcr.io\  
--docker-username='${oauthtoken}'\  
--docker-password=$NGC_CLI_API_KEY
```

Create Kubernetes secret with the NGC API key to launch NIMs

```
kubectcl create secret generic ngc-nim-secret  
--from-literal=NGC_CLI_API_KEY=$NGC_CLI_API_KEY
```

Install NVIDIA NGC CLI

```
sudo apt install unzip  
  
wget --content-disposition  
https://api.ngc.nvidia.com/v2/resources/nvidia/ngc-apps/ngc_cli/versions/3.43.0/files/ngccli_linux.zip -O ngccli_linux.zip && unzip ngccli_linux.zip  
  
wget --content-disposition  
https://api.ngc.nvidia.com/v2/resources/nvidia/ngc-apps/ngc_cli/versions/3.52.0/files/ngccli_linux.zip -O ngccli_linux.zip && sudo apt install unzip && unzip ngccli_linux.zip  
  
echo "export PATH=\"$PATH:$(pwd)/ngc-cli\"" >> ~/.bash_profile && source ~/.bash_profile
```

Configure NGC CLI client. Enter API key, enter org, leave everything else as default

```
ngc config set
```

Download the model

```
ngc registry model download-version  
"mphexwv2ysej/meta-llama3-8b-instruct:0515-db4a5074-trtllm10-1xa100-fp16"  
  
ngc registry model download-version "nvidia/llama3-8b-instruct:1.0"  
  
sudo mkdir /mnt/nim  
  
sudo mv meta-llama3-8b-instruct_v0515-db4a5074-trtllm10-1xa100-fp16 /mnt/nim/
```

```
sudo mv llama3-8b-instruct_v1.0 /mnt/nim/
```

~~# Enable the NodeSelector feature of KServe to allow a NIM to request different GPU types~~

```
kubectl patch configmap config-features -n knative-serving --type merge -p  
'{"data":{"kubernetes.podspec.nodeselector":"enabled"}}'
```

Create a PVC called nim-pvc in the cluster

```
cat > nim-model-volume.yaml << EOL  
  
apiVersion: v1  
  
kind: PersistentVolume  
  
metadata:  
  name: nim-pv  
  
spec:  
  capacity:  
    storage: 100Gi  
  
  volumeMode: Filesystem  
  
  accessModes:  
    - ReadWriteMany  
  
  persistentVolumeReclaimPolicy: Retain  
  
  storageClassName: microk8s-hostpath  
  
  local:  
    path: /mnt/nim  
  
  nodeAffinity:  
    required:  
      nodeSelectorTerms:  
        - matchExpressions:  
          - key: kubernetes.io/hostname  
            operator: In
```

```
    values:
      - `hostname`

---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nim-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: microk8s-hostpath
resources:
  requests:
    storage: 100Gi
EOL
```

```
kubectl create -f nim-model-volume.yaml
```

Create ClusterServingRuntime

```
cat > cluster-serving-runtime.yaml << EOL
apiVersion: serving.kserve.io/v1alpha1
kind: ClusterServingRuntime
metadata:
  name: nim-meta-llama3-8b-instruct-24.05.rc11
spec:
  annotations:
    prometheus.kserve.io/path: /metrics
```

prometheus.kserve.io/port: "8000"

serving.kserve.io/enable-metric-aggregation: "true"

serving.kserve.io/enable-prometheus-scraping: "true"

containers:

- env:

- name: NIM_CACHE_PATH

value: /tmp

- name: NGC_API_KEY

valueFrom:

secretKeyRef:

name: ngc-nim-secret

key: NGC_CLI_API_KEY

~~image: nvcr.io/mpheerwv2ysej/meta-llama3-8b-instruct:24.05.rc11~~

image: nvcr.io/nim/meta/llama3-8b-instruct:1.0.0

name: kserve-container

ports:

- containerPort: 8000

protocol: TCP

resources:

limits:

cpu: "1"

memory: 16Gi

requests:

cpu: "1"

memory: 16Gi

volumeMounts:

- mountPath: /dev/shm

```
  name: dshm
imagePullSecrets:
- name: ngc-docker-secret
protocolVersions:
- v2
- grpc-v2
supportedModelFormats:
- autoSelect: true
  name: nim-meta-llama3-8b-instruct
  priority: 1
  version: "24.05"
volumes:
- emptyDir:
    medium: Memory
    sizeLimit: 16Gi
  name: dshm
EOL
```

```
kubectl create -f cluster-serving-runtime.yaml
```

Create InferenceService

```
cat > inference-service.yaml << EOL
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  annotations:
    autoscaling.knative.dev/target: "10"
```



```
name: nim-meta-llama3-8b-instruct-1xgpu
spec:
  predictor:
    minReplicas: 1
  model:
    modelFormat:
      name: nim-meta-llama3-8b-instruct
  resources:
    limits:
      nvidia.com/gpu: "1"
    requests:
      nvidia.com/gpu: "1"
  runtime: nim-meta-llama3-8b-instruct-24.05.rc11
  storageUri: pvc://nim-pvc/
EOL
```

```
kubectl create -f inference-service.yaml
```

Watch the status of the pod created until it becomes ready

```
kubectl describe pod nim-meta-llama3-8b-instruct-1xgpu
```

Note: It can take a while for the pod to become functional. The warning message "Readiness probe failed" can be ignored. Continue with the instructions below and start testing queries.

Get the IP address of the private predictor

```
KSERVE=$(kubectl get svc |grep private |awk '{print $3}')
```

Validate that the NIM is running by posting a query against the KServe endpoint

```
curl http://$KSERVE/v1/chat/completions -H "Content-Type: application/json" -d '{
```

```
"model": "meta-llama3-8b-instruct",  
"messages": [{"role": "user", "content": "What is KServe?"}]  
'
```

```
curl http://$KSERVE/v1/chat/completions -H "Content-Type: application/json" -d '{  
"model": "llama3-8b-instruct",  
"messages": [{"role": "user", "content": "What is KServe?"}]  
'
```

Cleanup

```
sudo snap remove --purge juju  
sudo snap remove --purge microk8s  
sudo rm -rf /home/$USER/.local/share/juju  
rm -rf ~/.kube/
```

Gokhan Cetinkaya
Alliances Engineer
gokhan.cetinkaya@canonical.com
www.linkedin.com/in/gcetinkaya

References:

<https://github.com/NVIDIA/nim-deploy/tree/main/kserve>